

Memory-assisted prompt editing to improve GPT-3 after deployment

Aman Madaan^{*}, Niket Tandon^{*†}, Peter Clark[†], Yiming Yang

Language Technologies Institute, Carnegie Mellon University, Pittsburgh, PA, USA

[†] Allen Institute for Artificial Intelligence, Seattle, WA, USA

{amadaan, yiming}@cs.cmu.edu

{nikett, peterc}@allenai.org

Abstract

Large LMs such as GPT-3 are powerful but not immune to mistakes. For example, GPT-3 would mistakenly interpret "What word is similar to *good*?" to mean a homonym, while the user intended a synonym. Our goal in this paper is to effectively correct such errors via user interactions with the system but without retraining, which will be prohibitively costly. We pair GPT-3 with a growing memory of recorded cases where the model misunderstood the user's intents, along with user feedback for clarification or instruction. Such a memory allows our system to produce enhanced prompts for any new query based on the user feedback for error correction on similar cases in the past. On four tasks (two basic lexical tasks, two advanced ethical reasoning tasks), we show how a (simulated) user can interactively teach a deployed GPT-3, substantially increasing its accuracy over the queries with different kinds of misunderstandings by the GPT-3. Our approach is a step towards the low-cost utility enhancement for very large pre-trained LMs.¹

1 Introduction

GPT-3, while powerful, is not immune to mistakes (Marcus, 2021). The typical remedy of retraining with more data is not easy for these huge models, due to the prohibitive cost and infrastructure requirements. In such cases, even if users observe the model making a mistake repeatedly, there are no avenues to provide feedback to the model.

One failure mode is misinterpreting a user's instruction, or *intent*. For example, in Figure 1 the user has asked for a synonym, but the model misinterpreted the request as asking for a homonym. Depending on the user's expertise, and dialect, tasks may be expressed in various ways, leading to potential misunderstandings when the model encounters a new dialect or a poorly worded task.

^{*}Equal Contribution

¹Code and data is available at <https://github.com/madaan/memprompt>

Our memory enhanced GPT-3 implementation.

User: What word is similar to *good*?

GPT-3: The homonym of good is: wood.

User: "Similar to" means "with a similar meaning".

GPT-3: Noted [*writes to memory*]

User: What word is similar to *surprised*?

GPT-3: [*Retrieves and adds to prompt "Similar to" means "with a similar meaning"*].

The synonym of surprised is: amazed.

Figure 1: This paper enhances GPT-3 performance by looking up questions with a similar intent that received any user feedback. Our approach is simple because only the prompt needs to be updated with the retrieved relevant feedback, and no retraining is necessary.

Our goal is to allow users to correct such errors directly through interaction, and without retraining. Our approach is to pair GPT-3 with a growing memory of cases where the model misunderstood user's intent and was provided with corrective feedback. We use the feedback to clarify the intent of new questions. This feedback is question dependent, and thus the prompt for each sample is *edited* to adapt to the input. In this sense, our work can be seen as an instance of prompt engineering (Liu et al., 2021c) which involves editing the prompts. Our work adds interactivity to prompt engineering as it involves dynamically updating the prompt for every instance.

For systems deployed in the wild, the end users cannot be expected to give feedback on the correct output for a given input (e.g., they might not know how to solve the task in realistic settings, such as translation). To circumvent this, we have GPT-3 verbalize its *understanding* *u* of the user's intent (in addition to providing an answer), a skill learned using few-shot examples in the prompt. Now, the user can see how the model interpreted their instructions, and provide corrective feedback *fb* if that interpretation was incorrect.

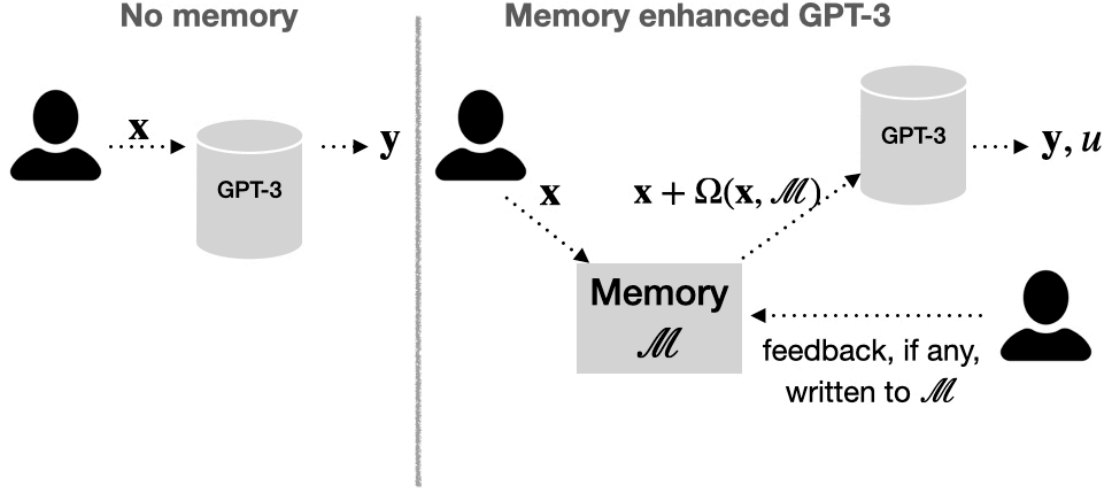


Figure 2: Proposed architecture: (left) GPT-3 does not account for user feedback. (right) MEM-PROMPT maintains a memory \mathcal{M} of corrective feedback, and searches for feedback from prior queries with a similar intent as x using a retrieval function Ω . x is then concatenated to the retrieved feedback and appended to the prompt for querying GPT-3. Users can also give new feedback on the model’s task understanding u , then added to \mathcal{M} .

Figure 1 presents a sample interaction between a user and GPT-3 that our setup enables. The model was asked for a similar word. However, the model’s (incorrect) task understanding u was “The homonym of good is”. The user can detect such discrepancy between the intended and interpreted task instruction, and can provide feedback fb as “*similar to means with a similar meaning*”, clarifying that they actually wanted a synonym. Crucially, note that such instructional correction is feasible *even if the user does not know the correct answer to their question*, as they are critiquing the model’s understanding of their intent, rather than the answers themselves. Thus, our setup **does not require** the users to be experts at tasks being solved, another advantage of our approach.

Further, it is desirable to have a system that can leverage past feedback on new, unseen examples for prompt-editing. To this end, we maintain a memory \mathcal{M} of such feedback as a set of key-value pairs, where the key is a misunderstood question, and the value is the user’s feedback to correct that misunderstanding. Given a new question, we check if the model has made a mistake on a similar question earlier, by querying the memory for a similar question. If found, append the corresponding feedback to the question prompt. This mechanism aims to prevent the model from making the same type of mistake twice. This failure-driven reminding mechanism draws inspiration from the theory of recursive reminding in psychology (Jacoby and

Wahlheim, 2013), which suggests humans index error corrections in the context in which those errors occurred.

This paper sets out the general architecture, along with simple, proof-of-concept implementations of its components. We then demonstrate the system on four tasks, using simulated user feedback: (1) lexical relations (e.g., antonyms, Figure 1), (2) word scrambling (e.g., anagrams), (3) ethics (with user feedback being the appropriate *class* of ethical consideration, e.g., “it is about cheating”, using a small set of categories), and (4) ethics (with user feedback being natural language). We find that in all cases, GPT-3’s accuracy significantly increases with time, without retraining, as it is able to use corrective feedback from earlier examples to avoid similar misunderstandings on future examples. Our contributions are thus a general architecture and illustrative implementation showing how user feedback might continuously improve model performance, without retraining, in a few-shot prompt setting.

2 Related work

Our method builds upon the recent advances in prompt-tuning and few-shot prompting.

Our use of recalled memories is a form of “prompt engineering”, where GPT-3’s behavior is modified by adding to the query (prompt) (Le Scao and Rush, 2021). Like others, we use GPT-3 with *few-shot prompting*, where the prompt consists of a

prefix *prefix* containing a few input-output “training” examples of the task, followed by the **input** x , e.g., a question, to operate on. However, while prior work has focused on constructing better prefixes, e.g., dynamically selecting good “training” examples based on the question (Liu et al., 2021a), or even representing the prefix latently (Li and Liang, 2021), our work elaborates the input x itself to clarify the intended task, by adding user feedback fb from previous misunderstandings.

Similarly, our work can be seen as a form of retrieval-augmented QA. Extensive prior work has used retrievals from a text corpus to aid QA, e.g., (Pan et al., 2019; Guu et al., 2020), or retrievals of prior QA pairs for nearest-neighbor QA (Khandelwal et al., 2020). In contrast, we are retrieving from a dynamic memory of user feedbacks.

The idea of failure-driven reminding and dynamic memory date back several decades, e.g., (Schank, 1983; Riesbeck, 1981). Our work resurrects these ideas in a modern context.

Learning from instruction has also become important for large LMs, where models can perform a task based on direct instruction rather than examples (Wei et al., 2021; Mishra et al., 2021). Our work extends this by adding an adaptive component for when those instructions are misinterpreted. An additional contribution of our work is in proposing a novel strategy for taking feedback from the users in cases where the model misunderstands an instruction. Specifically, as part of the output, the model repeats its understanding of the instruction. While it may not be possible for an end user to provide meaningful feedback on the output itself, giving feedback on the understanding of the instruction is more feasible.

Given an erroneous answer, our approach aims to modify the model’s behavior through prompting. An alternative, recently explored approach is “model editing” - updating the model itself by modifying its parameters to fix erroneous answers (Mitchell et al., 2021; De Cao et al., 2021; Hase et al., 2021). However, model editing approaches have to date only been demonstrated in a limited context (e.g., correcting a single error), and even then can lead to uncontrollable out of scope changes (Mitchell et al., 2021). In contrast, our goal is not just to correct a specific prediction, but to generalize that correction for new problems by collecting feedback to clarify the misunderstanding and without risking damage to the model’s basic

problem-solving acumen.

Finally, our work is a simple example of debugging and learning via dialog. While system debugging through dialog has been explored in many contexts, e.g., (Hixon et al., 2015; Wang et al., 2016; Davis, 1977), our novel contribution is dialog about the model’s understanding of the user’s intent.

3 Approach

3.1 Memory enhanced GPT-3 architecture

In our setup, given an input x , a model generates an output y and a sentence u expressing its understanding of the task, a skill learned through few-shot examples in the prompt (Appendix B). The user can then critique u by providing natural language feedback fb . This is feasible even if the user does not know the correctness of y because they are critiquing the *model’s understanding of their intent* rather than the answers themselves.

Given a new query, MEM-PROMPT uses fb from similar, prior queries to enrich the (few-shot) prompt p . We use the principle that if x_i and x_j have similar errors (i.e., $x_i \sim x_j$), then their feedbacks fb_i and fb_j should be exchangeable ($x_i \sim x_j \Leftrightarrow fb_i \sim fb_j$). Fig. 2 gives an overview of MEM-PROMPT, with the following components:

Memory \mathcal{M} : \mathcal{M} is a growing table of key (x_i) - value (fb_i) pairs that supports read, write, and lookup operations. The write operation is used whenever a user gives new feedback.

Lookup $\Omega(x, \mathcal{M})$: Ω is a learned retriever that matches the query= x against all the keys of \mathcal{M} .

Combiner $\mathcal{C}(x, \Omega(x, \mathcal{M}))$: A gating function allowing irrelevant, retrieved feedback to be ignored.

Prompter $\mathcal{P}(p, \mathcal{C})$: \mathcal{P} passes the output of \mathcal{C} to GPT-3 prompt.

Few-shot prompting Let us briefly recap few-shot prompting with GPT-3. Consider a general setup where given an input x , a model is expected to generate an output y . In a few-shot prompting mode (Brown et al., 2020), a prompt p consists of k (x, y) “in-context” examples, i.e., $p = x_1.y_1 \# x_2.y_2 \dots \# x_k.y_k$, where $\#$ is a token separating examples. During inference, the user inputs a question x_i , and the model is fed $p \# x_i$ (i.e., the question suffixed to the prompt) and is expected to generate the answer y_i as a continuation.

Task (fb type)	($\mathbf{x} \rightarrow \mathbf{y}$)	u and fb
Lexical relations (INS)	x : What sounds like good? y : wood	u : Question is asking for a synonym. fb : No, I want a homonym.
Word scrambling (INS)	x : Find the right word given this cycled word: elylarg y : largely	u : The question is about anagram. fb : No, its about uncycling a word.
Ethical reasoning (CAT)	x : Turning my blender on at 3AM y : It's bad.	u : Question is about authority. fb : No, it is about harm.
Ethical reasoning (NL)	x : John has started using again after his mother passed y : It's bad.	u : Question is about spending money. fb : No, it is about drug use.

Table 1: Feedback types and demonstration of understanding: our system leverages user feedback to prevent failures caused due to a misunderstanding of the task (INS) or semantics of the input (CAT and NL). We achieve this by having the model articulate an understanding **u**, on which a user can provide feedback using **fb**.

\mathcal{P} supplements this few-shot prompting workflow, with a memory of user feedbacks from $\mathcal{C}()$. To enable the model to react to such feedback, we include k samples of the form $(\mathbf{x}, \mathbf{fb} \rightarrow \mathbf{u}, \mathbf{y})$ in the prompt, so the question contains **fb**.

3.2 Feedback on model’s understanding

In the setup $(\mathbf{x} \rightarrow \mathbf{u}, \mathbf{y})$, there are three modes of failure for a model:

- Task instruction understanding: this is especially concerning in a multi-tasking setup, where the model may consider the question to be about a different task than the one user intended.
- Task nuanced understanding (error on **u**): this type of error occur when the model understands the task type, but misunderstands the subtle intent in a question.
- Task modeling: if the task is clearly understood, but the answer is not correct, then it requires updating the model parameters, see Section §2 for related work on model editing. However, existing approaches do not scale to very large LMs such as GPT-3.

The first two failure modes are due to the inability of the model to understand the input, and are our focus for this work. This paper provides an architecture for a user to critique on model failures.

Both these types of understanding (task understanding and task intent understanding) can be critiqued via feedback. In the case of NL feedback, the user may provide a counter argument refuting **u** as the possible intent. Similarly, in the case of categorical feedback, the user can critique the category provided by the model as its understanding of the situation. While feedback on the model output is our primary goal, we also experiment with settings where an Oracle is available to provide feedback on the labels (Section §4.2).

The model reacts to the feedback because some in-context samples are of the form: $(\mathbf{x}, \mathbf{fb} \rightarrow \mathbf{u}, \mathbf{y})$ and $(\mathbf{x} \rightarrow \mathbf{u}, \mathbf{y})$. We consider a diverse set of tasks $(\mathbf{x} \rightarrow \mathbf{y})$, **fb** and **u**, summarized in Table 1 and described next.

3.3 Tasks

We apply our approach to four tasks: (1) lexical relations (e.g., antonyms, Figure 1), (2) word scrambling (e.g., anagrams), (3) ethics (with user feedback being the appropriate *class* of ethical consideration, and (4) ethics (with user feedback being natural language). In each case, we simulate user feedback using additional data provided in the datasets. We describe these now.

3.3.1 Lexical Relations

The lexical relation task is to predict a word with a given lexical relationship to an input word. We use five relationships: synonym (*syn*), antonym (*ant*), homonym (*hom*, for our experiments, we define homonyms to be the set of words that have different spellings but identical pronunciation, like *ring* and *wring*), definition (*defn*), and sentence usage generation (*sent*). The sources for this dataset are listed in Appendix §C lists the sources for the dataset, and we release the dataset on the repository for this work.

We chose this task as each question can be asked in multiple ways (e.g., for synonym generation, the users might ask questions of the form *what is like*, *what has a similar sense*, *what is akin to*, *what is something like*, etc.) For each task, the prompt contains a few different variations, e.g., the variations for the homonym task include “*what is the homonym of <word>?*”, “*what sounds like <word>?*”. We create a dataset of $(\mathbf{x}, \mathbf{fb} \rightarrow \mathbf{u}, \mathbf{y})$ tuples using sentence templates, where **fb** clarifies the task in **x**. We then experiment in a **simulated** conversational setting, in which a user can ask the

model x (covering any of these five tasks). If the model gives the wrong answer to a query x , then fb is used as the simulated corrective feedback to the model.

3.3.2 Word Scrambling

For this task, given a word with its characters transformed, the model is expected to recover the original characters. There are four transformation operations the user can request: reversal of words (*rev*, *yppup* \rightarrow *puppy*), cycle letters in word (*cyc*, *atc* \rightarrow *cat*), random insertions (*rand*, *c!r ic/ke!t* \rightarrow *cricket*), and anagrams by changing all but the first and last (*anag1*, *eelhpnat* \rightarrow *elephant*) or all but the first and last 2 characters (*anag2*, *elapehnt* \rightarrow *elephant*). We use the original dataset provided by Brown et al. (2020).²

In the same way as for the lexical relations task, we specify the task description x using different phrasings, e.g., “rearrange the letters” (which the system sometimes misunderstands), and the (simulated) user feedback fb is a clearer task description, e.g., “The anagram is”. The system thus accumulates a set of x - fb pairs in memory after each failure, helping it avoid future misunderstandings of x through feedback retrieval.

3.3.3 Ethical Reasoning (2 tasks)

For ethical reasoning, we consider a setup where given a situation (e.g., *cheating on your partner*), the model is expected to provide a judgment on whether the situation is ethical or not (e.g., *it’s not okay*). In addition to providing a judgment on the ethics of the situation, the model also elucidates its understanding of what the question is about (e.g., *being loyal*). While the user may not know the answer, we posit that they would be able to provide feedback on the broader context. For example, if the model generates *being financially savvy* instead of *being loyal*, a user can still point out this problem and provide feedback.

As mentioned, our setup involves an input question (x), for which the model generates an answer (y) and an understanding of the question (u). The user then provides a feedback fb on model’s understanding, which is store to a memory. For subsequent queries, the system first retrieves a relevant feedback from the memory, and the feedback is combined with the question for the retrieval.

We use a subset³ of the dataset provided by DELPHI (Jiang et al., 2021). We simulate two different kinds of user feedback, using two of the annotations attached to each example in the Delphi dataset:

- **Categorical feedback (ERT-CAT):** In this setting, the model generates its understanding u of the situation by selecting one of 10 different possible categories of morality to which the situation might belong: care, loyalty, authority, fairness, sanctity, degradation, cheating, subversion, betrayal, and harm. (These categories are explicitly provided for each example in the Delphi dataset). The model is “taught” to generate a category u (as well as the okay/not-okay answer y to the ethical question) by being given a few examples in the prompt prefix, thus articulating which moral category it thinks is applicable. The simulated feedback fb is the gold category associated with the example in the question, if GPT-3 gets the answer wrong.
- **Natural language feedback (ERT-NL):** For this, we use the associated “rule of thumb” (RoT) annotation - a freeform general moral principle - attached to each example in the Delphi dataset. As before, the model is “taught” to generate its understanding u of the appropriate RoT (as well as the okay/not-okay answer y) using few-shot examples in the prompt prefix, and the gold RoT is used as (simulated) user feedback fb if GPT-3 answers the question incorrectly. To compile a challenging subset of the data for ERT-NL, we sample by input length, preferring long x , with a short feedback fb . Specifically, we use the top 1% of the inputs by length to create a challenging set of input situations (x). User feedback is then a similar rephrasing of the input.

We selected these tasks because situations that involve reasoning about similar ethical principles can utilize similar past feedback. For example, *sharing an extra umbrella with your friend if they don’t have one*, and *donating surplus food to the homeless* both involve *compassion*.

3.4 A Proof of Concept Implementation

Implementation of memory \mathcal{M} We implement \mathcal{M} using x as the key and the corresponding feedback fb as value. Given a question x_i , if the user detects that the model has misunderstood the question, they may provide a fb_i with probability $Pr(f_i)$. The feedback is stored in a memory \mathcal{M} , with x_i as

²<https://github.com/openai/gpt-3/tree/master/data>

³social norms dataset (social-chemistry-101) <https://github.com/mbforbes/social-chemistry-101>

Question	Feedback
A word pronounced as fellow ?	I want a word that sounds similar!
What is dissimilar to delicious ?	Give me the reverse of delicious
What is a word like great ?	Wrong! I want something similar ✓
How do I use melancholy ?	No...I wanted a sample sentence
What is on the lines of pretty ?	I was looking for a similar word
Could you expand on browser ?	I actually wanted a definition

1. Query the memory

↑

q : What is akin to quick ?

2. Retrieve relevant feedback

↓

fb: Wrong! when I mention like, I want something similar

Figure 3: Sample snapshot of memory for lexical QA.

the key and fb_i as the value. For a subsequent question \mathbf{x}_j , the retriever Ω (described below) checks if a similar question appears in memory. If yes, then the corresponding feedback is attached with the question and fed to the model for generation.

For example, the model might misunderstand a question asking for synonym, e.g., *what is akin to fast ?* as one that requires antonyms. As mentioned, in our setup, the model generates its understanding of the task \mathbf{u} , and not just the answer to the question. The user, by inspecting $\mathbf{u} = \text{The opposite of fast is:}$ might determine that the model has misunderstood them, and give feedback *i wanted a synonym*, which gets stored in \mathcal{M} . If a similar question (e.g., *what is akin to pretty ?*) is asked later by the same or a different user, the corresponding feedback (*i wanted a synonym*) is attached with the question to generate the answer. Figure 3 illustrates a sample memory for this task.

Implementation of retriever Ω An incorrect feedback might cause the model to make a mistake, thus necessitating a good retrieval function. In our setting, we use two different retrieval functions:

(1) Semantic similarity: the query is encoded using Sentence transformers (Reimers and Gurevych, 2019), and we use cosine distance with a threshold of 0.9 to find a matching key \mathbf{x}_m .

(2) Lexical similarity: We also experiment with low-resource settings for which trained retrieval is not an option. In such cases, we rely on heuristics for similarity matching (details in Appendix §D).

Implementation of combiner \mathcal{C} \mathcal{C} concatenates x and fb retrieved by Ω . We leave space for future work to do gating in a more principled manner.

Implementation of prompter \mathcal{P} \mathcal{P} concatenates \mathcal{C} at the end of p . Future work can employ strategies in recent literature on prompt-fine tuning (Zhao et al., 2021) to best combine fb with p e.g., deciding the position of p or format of \mathcal{C} ’s output for best gains.

Although the model has not changed, adding fb corrects its erroneous behavior because we provide a few positive “training” examples containing feedback ($\mathbf{x}, \text{fb} \rightarrow \mathbf{u}, \mathbf{y}$) in the prompt (Appendix B).

4 Experiments

Baselines We compare our system, MEM-PROMPT (memory-assisted prompt editing) with two different baselines:

1. **NO-MEM** This is the standard GPT-3⁴ in few-shot prompting mode, with the suggested parameters (Appendix §A). Input is $\mathbf{p} \# \mathbf{x}_i$ (i.e., question \mathbf{x}_i appended to prompt \mathbf{p}). It generates answer \mathbf{y}_i and its understanding of the user’s intent \mathbf{u}_i .
2. **GROW-PROMPT**: Similar to NO-MEM, but the \mathbf{p} is continuously grown with a subset of memory \mathcal{M} that can fit within the prompt (max. 2048 tokens). The most recent subset of \mathcal{M} of memory inserted is inserted in the prompt. The ethical reasoning tasks (ERT) involve long examples, and the initial prompt itself takes close to the max allowed tokens. Thus, the GROW-PROMPT setup is only provided for the lexical relations and word scrambling tasks.

Metrics We use two different metrics:

- $\text{Acc}(\mathbf{y})$: % of cases where answer matched the ground truth.
- $\text{Acc}(\mathbf{u})$: % of cases where the model’s understanding of user’s intent is correct. As discussed in Section 3.2, depending on the task, the model generates its understanding on either the instruction or semantics of the question.

4.1 Main result: MEM-PROMPT improves GPT-3 accuracy

The most important research question is whether pairing GPT-3 with MEM-PROMPT improves performance. We validate this using a diverse set of tasks.

⁴We use GPT-3-175B (davinci) for all experiments.

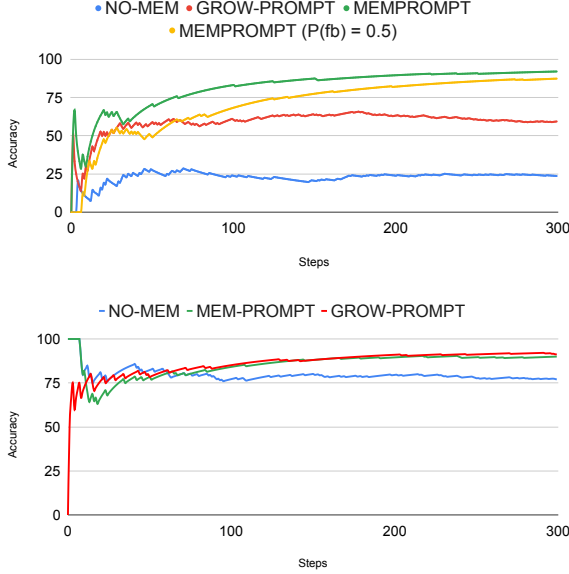


Figure 4: **Main result** Avg. performance over 300 data points on five lexical tasks (top) and word scramble tasks (bottom). The baselines are NO-MEM (in blue), which is GPT-3 in a few-shot setup. GROW-PROMPT (in red) appends the prompt with memory. MEM-PROMPT is shown in green. The horizontal axis records the number of steps (or equivalently time). For both the memory assisted systems (GROW-PROMPT and GROW-PROMPT), the accuracy increases with time as memory is filled up with feedback from past errors.

4.1.1 Ethical reasoning tasks

Table 5 presents results from running MEM-PROMPT on the DELPHI dataset (1,000 points in the test set). Recall from §3.3 that there are two different kinds of feedback that we provide on DELPHI questions: CAT and NL feedback. MEM-PROMPT gets over 10% relative improvement in ERT-NL and over 25% relative improvement in ERT-CAT.

Feedback improves accuracy over time Figure 6 demonstrates that the instruction accuracy increases over time. We also see that if the user frequently provides clarifications then the model gets better more quickly. Figure 5 shows the label accuracy also improves over time. The baseline (NO-MEM) quickly saturates after 200 timesteps, but MEM-PROMPT continues to improve. This is one of the advantages of engaging humans in the loop - the model keeps improving with time. It is also evident from these charts that as instruction accuracy improves, the label accuracy also improves.

Error analysis: Ethical-NL In both the ERT NL and CAT tasks, one of the primary source of label errors is confusion between labels such as OKAY

and GOOD because of the nuanced differences e.g., input = teaching your child a musical instrument. MEM-PROMPT predicted GOOD, while the expected answer was OKAY. Similar trends in this dataset were also observed by Jiang et al. (2021)

We randomly sampled from the ERT-NL test set where the model generates an incorrect understanding (i.e., $Acc(\mathbf{u}) = 0$ based on exact match). Our goal is to understand the typical errors made by the model and use the analysis to calibrate the findings in Table 5. We select ERT-NL for the analysis because it involves free-form natural language which is difficult to study quantitatively.

- **Correct, lexically variant understanding (30%):** Exact match underestimates the performance of our model (as the task involves generation). About 30% of the model understanding is a lexical variation of the reference gold understanding. For example, *telling a spouse your true feeling* vs. *loving your partner*. Notably, we find that the generated label in these cases is still correct. (Example in Table 2, row 1)
- **Distracted understanding (50%):** Majority of the errors in understanding (and the label) come from the model being distracted by an unimportant detail. Retrieval issues account for about 30% of the errors within this category (e.g., a partially related situation in the memory where the expected understanding does not transfer to the query.) (See Table 2, row 2)
- **Retrieval failures (18%):** These errors are caused by an irrelevant retrieved understanding from the memory. We believe that a better retrieval function (e.g., one that models analogies between input situations) can potentially help alleviate these issues, and we leave such exploration for future work. (See Table 2, row 3)

Canonical examples of these error categories are shown in Table 2. We find that with time, more relevant past examples are fetched. Table 7 shows some examples of retrieved examples over time.

4.1.2 Word Reasoning Tasks

For these tasks, we compare gold \mathbf{u}^* and generated \mathbf{u} based on some hard-coded linguistic variations (e.g., *the antonym is matches the opposite is*). Failure to generate \mathbf{u} is also considered incorrect. While we do not explicitly evaluate the accuracy of the task, we found a near-perfect correlation between the accuracy of \mathbf{y} and \mathbf{u} (i.e., if the GPT-

Error category	Input	Output (MEM-PROMPT)	Reference
Correct, lexically variant understanding	when I said "i do" to my wife, I am now realizing I only meant "i'll try". does she deserve to know and a chance to back out?	This question is about: telling a spouse your true feelings.	This question is about: loving your partner.
Distracted understanding	i am 20 years old, lately I have been forgetting simple spellings while taking notes in class . I often bump into random objects around hurting myself.	This question is about: wanting to kill yourself. The answer is BAD.	This question is about: seeking medical attention if you feel there is something wrong with you. The answer is GOOD.
Retrieval failures	checked out my high school ex girlfriends instagram for the first time in two years after graduating. I'm happy I did. she's not as pretty as she used to be and I feel great.	[retrieved feedback] This question is about: looking at someone's social media. The answer is GOOD.	This question is about: judging people based on physical attributes. The answer is BAD.

Table 2: ERT NL task- error categories

3 understands the task correctly, the output was almost always correct).

Figure 4 reports the overall performance on the five lexical tasks overall. The accuracy improves substantially within 300 examples when using memory (in yellow) vs. no memory (in blue). Table 3 breaks down the performance by tasks. We note again that we are operating in a few-shot prompting regime (i.e., there is no training data over which we train). The fact that the model saturates within 300 examples shows that our method can continue to improve. The performance of GROW-PROMPT (red) lies in between, showing that non-selective memory is partially helpful, although not as effective as failure-driven retrieval (our model). However, GROW-PROMPT is $\sim 3\times$ more expensive (larger prompts) and cannot scale beyond the 2048 tokens limit. Our model MEM-PROMPT substantially outperforms both the baselines, showing the effectiveness of failure-driven reminding. We also found that the retrieved feedback from memory was effective 97% of the time; only in $\approx 3\%$ of cases feedback had no positive effect.

We also note that the performance gains achieved by MEM-PROMPT are less dramatic for word-level tasks. This is explained by the fact that task descriptions for the word scrambling tasks are less ambiguous (Section §3.3), preventing the model from getting confused by users' instructions.

Persistent use of memory accelerates performance : When the memory is used for every example (green line in Fig 4, top), the performance improves quickly as compared to the yellow line, where **fb** from memory is drawn with $Pr(\mathbf{f}_i) = 0.5$.

model	syn	ant	hom	sent	defn	all
NO-MEM	0.58	0.43	0.13	0.30	0.39	0.37
GROW-PROMPT	0.71	0.87	0.75	0.92	0.76	0.80
MEM-PROMPT	0.99	0.98	0.98	0.98	0.96	0.98

Table 3: Results lexical QA tasks. Across all tasks, MEM-PROMPT has the best performance.

model	anag1	anag2	cyc	rand	rev	all
NO-MEM	0.81	0.47	0.95	0.98	0.62	0.77
GROW-PROMPT	0.86	0.89	0.93	0.96	0.90	0.91
MEM-PROMPT	0.81	0.83	0.98	0.95	0.93	0.90

Table 4: Results on word scramble QA tasks. Across all tasks, GROW-PROMPT and MEM-PROMPT outperform NO-MEM.

model	ERT-CAT	ERT-NL
NO-MEM	46.6	34.4
GROW-PROMPT	-	-
MEM-PROMPT	58.6	38.5

Table 5: Results on ERT. MEM-PROMPT outperform NO-MEM for both NL and CAT feedback. This is over 1000 test points each for ERT-CAT and ERT-NL. MEM-PROMPT improves over the NO-MEM setup for both categorical and the more challenging ERT-NL setup, which contains longer and ambiguous input situations.

%	Question (q)	Example 1	Example 2	Example 3
1	what highschool did harper lee go to?	what did st augustine do?	who is keyshia cole dad?	when did charles goodyear invented rubber?
4	what college did albert einstein go to?	what highschool did harper lee go to?	who did tim tebow play college football for?	what timezone is utah in?
40	where did john mayer go to college?	what school did michael jackson go to high school?	where did derek fisher go to college?	what style of music does john mayer play?
90	where did scott fitzgerald go to college?	who was f. scott fitzgerald?	where did otto frank go to college?	where did derek fisher go to college?

Table 6: Relevant examples fetched with time: as time proceeds, the examples fetched from the memory become increasingly relevant to the input question, leading to increasingly accurate predictions.

Timestep	Question	Neighbor
12	is not wanting a dirty person harassing you morally ok?	is not wanting to hear foul language morally good?
70	when adam decided to stop living a lie and announces he's gay, are you expected to pretend you're straight when you're gay?	are you expected to pretend you're straight when you're gay?
85	Should I help someone when they are having trouble?	must you intervene if you see someone taking advantage of another disabled person?
230	needing a hug to feel good.	is wanting to cuddle with loved ones okay?

Table 7: Relevant examples retrieved at increasing timesteps: as time proceeds, the examples fetched from the memory become relevant to the input question, leading to accurate predictions.

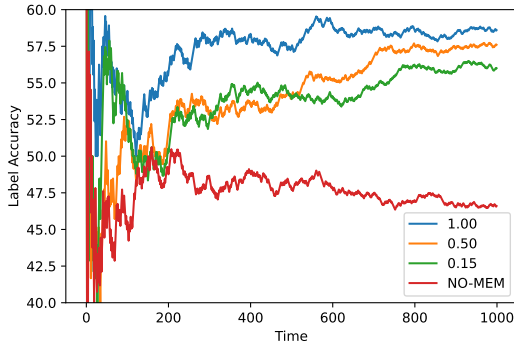


Figure 5: Label accuracy increases with time for ERT-CAT.

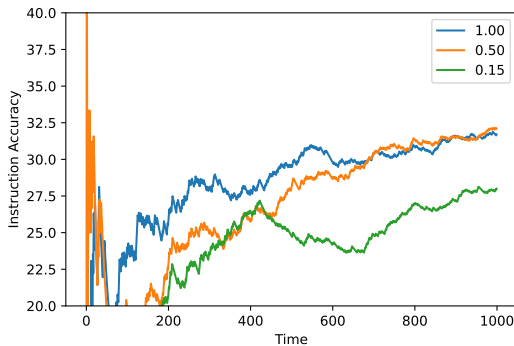


Figure 6: Instruction accuracy sharply increases with a larger clarification probability for ERT-CAT.

4.2 MEM-PROMPT with label feedback

Our current approach requires the model to verbalize its understanding of the question, on which a user provides feedback. Such a setup might not be possible, for instance, due to the nature of questions. Can MEM-PROMPT be effectively used in such settings as well? To investigate this, we experiment with factual question answering on the WEBQA dataset (Berant et al., 2013), and use the test set provided by Berant et al. (2013) for all experiments (2032 questions). The WEBQA dataset consists of factual questions (*which language is spoken in Canada?*) with multiple answers (*English, French*), and is a popular dataset for benchmarking the performance of GPT-3 on question answering in a few-context prompting setup.

Inference Let k be the number of examples (i.e., question-answer pairs) in the prompt. For a given question q , We keep half ($k/2$) examples fixed in the prompt, whereas the other half $k/2$ examples are retrieved from a memory of feedback M . As before, on receiving a question q , consults a memory M to see if a similar question has been asked before. However, different from earlier setups, in this case, we retrieve $k/2$ most similar questions from the memory M on **which the system has been wrong earlier**. The corresponding true answers are also retrieved. These $k/2$ retrieved question-answer pairs are combined with the $k/2$ fixed ques-

tions to create a prompt, and query GPT-3. Let a' be the generated answer.

Growing memory of errors M In our setup, we assume an expert user (or a teacher) that knows the true answer a for a given query q . The expert user compares the GPT-3 generated answer a' with a . If the generated answer is correct ($a' = a$), no further action is taken. If not, the entry $((q, a))$ is added to the memory M . As time passes, M is populated with an increasing number of challenging examples that the model has been wrong on. Thus, the retrieved $k/2$ examples get more relevant with time, aiding the accuracy. In the experiments, we set $k = 16$ due to budget constraints (note that the setups used in Liu et al. (2021b) and Brown et al. (2020) set $k = 64$, but their results are comparable to our baseline with $k = 16$).

Results Similar to ERT and word reasoning tasks, a memory of errors helps in increasing accuracy with time over 3,000 points in the test split of the WEBQA dataset (Figure 7). This is expected, as M gathers more examples on which GPT-3-175B has been wrong before. Adding these examples in the prompt avoids the model in repeating these mistakes.

To check if examples that belong to a similar domain improve with time, we cluster the questions in the test set of WEBQA, and randomly select three clusters for our analysis. Table 12 shows the top three of the 8 ($k = 16/2$) examples retrieved from M for the *alma mater* cluster.⁵ All of these questions relate to the alma mater of famous personalities. As the inference begins (with an empty M), the examples are not relevant to q . However, towards the end, almost all the samples are relevant to the given question.

5 Conclusion

We have presented a simple, novel, memory-enhanced GPT-3 that allows users to interact and improve the model without retraining. A key insight is to have the model articulate not just its answer but also its understanding of the user’s intent, providing an avenue for feedback. Our implementation of system components are illustrative, not definitive; rather, the goal of this paper is to suggest a general architecture for future researchers, where more sophisticated component implementations can be designed. This architecture is sig-

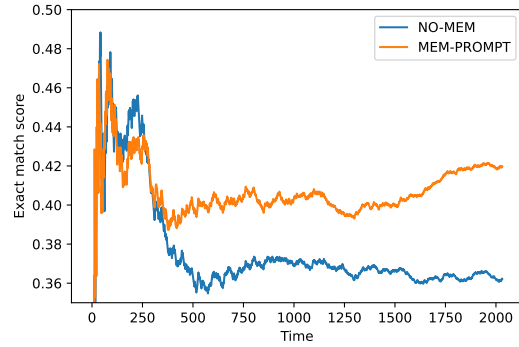


Figure 7: Instruction accuracy vs. time for WEBQA.

nificant as it suggests how deployed systems with fixed models can still be dynamically taught by interacting with end-users, potentially improving their performance and broadening their utility.

⁵Additional examples are included in Appendix §F.

References

- Jonathan Berant, Andrew Chou, Roy Frostig, and Percy Liang. 2013. Semantic parsing on freebase from question-answer pairs. In *Proceedings of the 2013 conference on empirical methods in natural language processing*, pages 1533–1544.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, and Others. 2020. [Language models are few-shot learners](#). In *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*.
- Randall Davis. 1977. Interactive transfer of expertise: Acquisition of new inference rules. *Artif. Intell.*, 12:121–157.
- Nicola De Cao, Wilker Aziz, and Ivan Titov. 2021. [Editing factual knowledge in language models](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6491–6506, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Kelvin Guu, Kenton Lee, Z. Tung, Panupong Pasupat, and Ming-Wei Chang. 2020. Realm: Retrieval-augmented language model pre-training. *ArXiv*, abs/2002.08909.
- Peter Hase, Mona Diab, Asli Celikyilmaz, Xian Li, Zornitsa Kozareva, Veselin Stoyanov, Mohit Bansal, and Srinivasan Iyer. 2021. Do language models have beliefs? methods for detecting, updating, and visualizing model beliefs. *arXiv preprint arXiv:2111.13654*.
- Ben Hixon, Peter Clark, and Hannaneh Hajishirzi. 2015. [Learning knowledge graphs for question answering through conversational dialog](#). In *Proceedings of the 2015 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 851–861, Denver, Colorado. Association for Computational Linguistics.
- Larry L. Jacoby and Christopher N. Wahlheim. 2013. On the importance of looking back: The role of recursive reminders in recency judgments and cued recall. *Memory & Cognition*, 41:625–637.
- Liwei Jiang, Jena D Hwang, Chandra Bhagavatula, Ronan Le Bras, Maxwell Forbes, Jon Borchardt, Jenny Liang, Oren Etzioni, Maarten Sap, and Yejin Choi. 2021. Delphi: Towards machine ethics and norms. *arXiv preprint arXiv:2110.07574*.
- Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2017. Billion-scale similarity search with gpus. *arXiv preprint arXiv:1702.08734*.
- Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. 2020. [Generalization through memorization: Nearest neighbor language models](#). In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net.
- Teven Le Scao and Alexander Rush. 2021. [How many data points is a prompt worth?](#) In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2627–2636, Online. Association for Computational Linguistics.
- Xiang Lisa Li and Percy Liang. 2021. [Prefix-tuning: Optimizing continuous prompts for generation](#). In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 4582–4597, Online. Association for Computational Linguistics.
- Bill Yuchen Lin, Wangchunshu Zhou, Ming Shen, Pei Zhou, Chandra Bhagavatula, Yejin Choi, and Xiang Ren. 2020. [CommonGen: A constrained text generation challenge for generative commonsense reasoning](#). In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 1823–1840, Online. Association for Computational Linguistics.
- Jiachang Liu, Dinghan Shen, Yizhe Zhang, Bill Dolan, Lawrence Carin, and Weizhu Chen. 2021a. What makes good in-context examples for gpt-3? *ArXiv*, abs/2101.06804.
- Jiachang Liu, Dinghan Shen, Yizhe Zhang, Bill Dolan, Lawrence Carin, and Weizhu Chen. 2021b. [What Makes Good In-Context Examples for GPT-3?](#) *arXiv:2101.06804 [cs]*. *ArXiv*: 2101.06804.
- Pengfei Liu, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2021c. Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing. *ArXiv*.
- Gary Marcus. [Experiments testing gpt-3’s ability at commonsense reasoning: results](#). [online]. 2021.
- Swaroop Mishra, Daniel Khashabi, Chitta Baral, and Hannaneh Hajishirzi. 2021. Natural instructions: Benchmarking generalization to new tasks from natural language instructions. *ArXiv*, abs/2104.08773.
- Eric Mitchell, Charles Lin, Antoine Bosselut, Chelsea Finn, and Christopher D. Manning. 2021. [Fast model editing at scale](#). *CoRR*.
- Kim Anh Nguyen, Sabine Schulte im Walde, and Ngoc Thang Vu. 2016. [Integrating distributional lexical contrast into word embeddings for antonym-synonym distinction](#). In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 454–459, Berlin, Germany. Association for Computational Linguistics.
- Xiaoman Pan, Kai Sun, Dian Yu, Jianshu Chen, Heng Ji, Claire Cardie, and Dong Yu. 2019. [Improving question answering with external knowledge](#). In *Proceedings of the 2nd Workshop on Machine Reading for Question Answering*, pages 27–37, Hong Kong, China. Association for Computational Linguistics.

- Nils Reimers and Iryna Gurevych. 2019. [Sentence-BERT: Sentence embeddings using Siamese BERT-networks](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992, Hong Kong, China. Association for Computational Linguistics.
- C. Riesbeck. 1981. Failure-driven reminding for incremental learning. In *IJCAI*.
- Roger Schank. 1983. *Dynamic Memory: A Theory of Reminding and Learning in Computers and People*. Cambridge University Press.
- Sida I. Wang, Percy Liang, and Christopher D. Manning. 2016. [Learning language games through interaction](#). In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2368–2378, Berlin, Germany. Association for Computational Linguistics.
- Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M. Dai, and Quoc V. Le. 2021. Finetuned language models are zero-shot learners. *ArXiv*, abs/2109.01652.
- Tony Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. 2021. Calibrate before use: Improving few-shot performance of language models. *ArXiv*, abs/2102.09690.

A Querying GPT-3-175B using OpenAI API

We use the OpenAI API for querying GPT-3-175B.⁶ The python code is listed below. Here, "PROMPT" is set to prompt shown in §B, followed by the input question x and feedback fb if applicable.

```
import os
import openai

openai.api_key = os.getenv("OPENAI_API_KEY")

response = openai.Completion.create(
    engine="davinci",
    prompt="PROMPT",
    temperature=0.7,
    max_tokens=64,
    top_p=1,
    frequency_penalty=0,
    presence_penalty=0
)
```

B Prompt

GPT3 is queried using a prompt p of example i/o behaviors, followed by the actual question x and (optionally) retrieved feedback fb . It then generates the understood intent u and answer y as a continuation. u and y are expressed a single sentence, e.g., "[The synonym for <word> is] [<word>]" Figure 8 shows this prompt p , containing a mixture of $(x \rightarrow u, y)$ and $(x, fb \rightarrow u, y)$ "training" tuples.

C Datasets for lexical question-answering tasks

As mentioned in Section §4, we focus on five different linguistic QA tasks. The source of data for each of these tasks is listed below:

1. The synonyms (syn) and antonyms (ant) were obtained from Nguyen et al. (2016).⁷
2. The homonyms (hom) were obtained using homz <https://github.com/cameronehrlich/homz>. We use the closest homonym returned by homz for each word in the English dictionary.
3. The definitions (defn) were sourced from *The Online Plain Text English Dictionary* <https://github.com/eddydn/DictionaryDatabase>

⁶<https://beta.openai.com/docs/introduction>

⁷<https://www.ims.uni-stuttgart.de/en/research/resources/experiment-data/lexical-contrast-dataset/>

4. Examples for usage in a sentence (sent) are from Commongen (Lin et al., 2020).

C.1 Templates

We manually created 15 task templates with three variants of phrasing the question for each task. Sample templates are shown in code listing 1. The data (word1, word2) in the code is initialized with the entries in the four sources mentioned above. The complete file is available in the anonymized code repository <https://anonymous.4open.science/r/memprompt-D548/templates.py>.

C.2 Sample questions

Tables 8, 9, and 9 list some sample x - y for settings where the question was asked as a linguistic variation, in Hindi, and in Punjabi, respectively.

D Finding similar questions in low-resource settings

In low-resource settings (e.g., queries in transcribed Punjabi or Hindi), we perform similarity matching between a given question and a question in the memory by using surface-form similarity. Specifically, we use Levenshtein distance to determine the closest query in the memory. We note that as the memory grows large, we can use mechanisms such as FAISS (Johnson et al., 2017) for trained memory, and suffix-trees for fast retrieval using surface form similarity.

We also experimented using queries in Hindi and Punjabi, with (English) feedback clarifying the queries' intent when GPT3 predictably misunderstands the task. Figure 10 confirms significant gains using memory in this OOV setting. This setup highlights the case when the user does not speak fluent English and uses mixed language code, e.g., transcription in English and mixing words from another language to ask questions.

E Sample results

Table 11 shows randomly sampled x - y pairs, and the corresponding y generated by GPT-3-175B and MEM-PROMPT. The complete set of outputs is located in the anonymized repository <https://anonymous.4open.science/r/memprompt-D548/results/results.csv>

What is the homonym for < wring > ?

#

the homonym for wring is ring END

#

how do you use < highway > in a sentence?

#

a sentence with highway is: soldiers stand guard along the [highway] END

#

can you define < camisole > ?

#

the definition of camisole is a sleeveless undergarment. END

#

What is the antonym for < prohibition > ?

#

the antonym for prohibition is permit END

#

What is the synonym for < surrogate > ?

#

the synonym for surrogate is substitute END

#

how do i use < fog > ? | clarification: when i ask for how do i use, i want a sentence.

#

a sentence with fog is: a rising sun burns the [fog] off a city END

#

What sounds like < sighted > ? | clarification: when I ask for sounds like, I want a homonym.

#

the homonym for sighted is cited END

#

what is like < provident > ? | clarification: when I ask for like, I want a synonym.

#

the synonym for provident is prudent END

#

can you define < rider > ? | clarification: when i ask for define, i want a definition.

#

the definition of rider is a person who is riding something. END

#

What is the opposite of < citation > ? | clarification: when I ask for opposite, I want an antonym.

#

the antonym for citation is award END

Figure 8: The prompt used for our tasks. During inference, an input question x_i , and optionally a feedback fb_i is appended after this prompt, and the model is expected to generate the answer y_i and its understanding of the question intent u_i as a continuation. The prompt contains examples of the form $(x \rightarrow u, y)$, expressed " $x \# u \ y$ END #", and $(x, fb \rightarrow u, y)$, expressed " $x \mid \text{clarification: } fb \# u \ y$ END #". (u and y are expressed together as a single sentence, e.g., "[The synonym for <word> is] [<word>].")

Find the right word after removing random letters from < t!r/e/a/s/u/r.e!s >

the word after removing symbols from t!r/e/a/s/u/r.e!s is treasures END

Find the original word after ignoring the punctuation and spaces in < e >

the word after removing symbols from e is elders END

Find the right word given this cycled word: < lprovisiona > ?

the uncycled version of lprovisiona is provisional END

Make a word while keeping the first and last char < vosiin > ?

the anagram 1 for vosiin is vision END

Find the original word that is interspersed in < f.i.n!e/p.i/x >

the word after removing symbols from f.i.n!e/p.i/x is finepix END

Find the right word given this rotated word: < cturalarchite > ?

the uncycled version of cturalarchite is architectural END

Find the original word after ignoring the punctuation and spaces in < s >

the word after removing symbols from s is straightforward END

Find the right word given this rotated word: < ibitioninh > ?

the uncycled version of ibitioninh is inhibition END

Figure out the word which has the same first two and the last two char < watsed > ? | clarification: when I want you to figure out the word which has the same first two and the last two char, I mean anagram 2.

the anagram 2 for watsed is wasted END

Make a word while keeping the first and last char < isucnase > ? | clarification: when I want you to make a word while keeping the first and last char, I mean anagram 1.

the anagram 1 for isucnase is issuance END

Unscramble everything except the first two and the last two char < acotrs > ? | clarification: when I want you to unscramble everything except the first two and the last two char, I mean anagram 2.

the anagram 2 for acotrs is actors END

Figure 9: The prompt used for our word scrambling tasks derived from [Brown et al. \(2020\)](#)

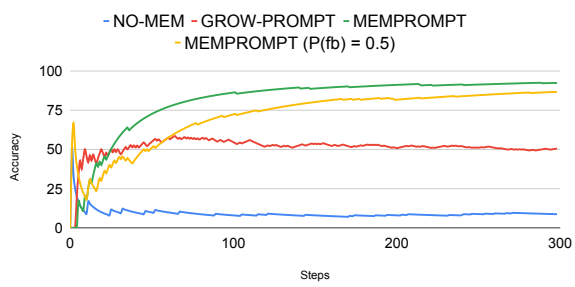


Figure 10: **Finding 2** Large gains on OOV setting by MEM-PROMPT over NO-MEM.

F Factual question answering

Tables 12 and 13 show additional examples for questions from WEBQA which get additionally relevant examples as time proceeds. The examples include questions that belong to the domains of Alma mater, Soccer, and Language.

```

1 templates = [
2     {
3         "type": "syn",
4         "template_id": "syn1",
5         "question": lambda word1: f"What is similar to < {word1} > ?",
6         "question_clarification": lambda word1: f"What is similar to < {word1} > ? |
clarification: when I ask for similar to , I want a synonym.",
7         "clarification": "clarification: when I ask for similar to , I want a synonym.",
8         "answer": lambda word1, word2: f"the synonym for {word1} is {word2}",
9     },
10    {
11        "type": "ant",
12        "template_id": "ant0",
13        "question": lambda word1: f"What is unlike < {word1} > ?",
14        "question_clarification": lambda word1: f"What is unlike < {word1} > ? |
clarification: when I ask for unlike , I want an antonym.",
15        "clarification": "clarification: when I ask for unlike , I want an antonym.",
16        "answer": lambda word1, word2: f"the antonym for {word1} is {word2}",
17    },
18    {
19        "type": "defn",
20        "template_id": "defn0",
21        "question": lambda word: f"< {word} > means what ?",
22        "question_clarification": lambda word: f"< {word} > means what ? | clarification:
when I ask for means what , I want a definition.",
23        "clarification": "clarification: when I ask for means what , I want a definition.",
24        "answer": lambda word, definition: f"the definition of {word} is {definition}",
25    },
26    {
27        "type": "sent",
28        "template_id": "sent1",
29        "question": lambda word: f"< {word} > can be used how ?",
30        "question_clarification": lambda word: f"< {word} > can be used how ? |
clarification: when I ask for can be used how , I want a sentence.",
31        "clarification": "clarification: when I ask for can be used how , I want a
sentence.",
32        "answer": lambda word, sentence: f"a sentence with {word} is: {sentence}",
33    }
]

```

Listing 1: "Sample templates for the five tasks."

Question (x)	Answer (y)	type
What is the opposite of < misconstrue > ?	the antonym for misconstrue is verify	ant
What is the opposite of < gross > ?	the antonym for gross is polite	ant
expand on < chelicera > ?	the definition of chelicera is One of the anterior pair of mouth organs	defn
what is a sentence that can be used to de- fine < mawseed > ?	the definition of mawseed is The seed of the opium poppy.	defn
what has a < bitt > like ring to it ?	the homonym for bitt is bit	hom
what can one confuse with < holed > ?	the homonym for holed is hold	hom
< spread > can be used how ?	a sentence with spread is: a couple of sheep are spread out in a field	sent
make something with < pot > ?	a sentence with pot is: bonsai tree in pot at zen garden .	sent
What is akin to < musician > ?	the synonym for musician is instrumental- ist	syn
What is akin to < zigzag > ?	the synonym for zigzag is move	syn

Table 8: Sample x-y pairs in English. The same type of question can be asked in multiple ways.

Question (x)	Answer (y)	type
< tabulate > ka ulta kya hai ?	the antonym for tabulate is randomize	ant
< foot > ka vilom kya hai ?	the antonym for foot is head	ant
< lettish > ka matlab kya hota hai ?	the definition of lettish is The language spoken by the Letts. See Lettic.	defn
< housing > ka arth kya hai ?	the definition of housing is An appendage to the hames or collar of a harness.	defn
sunne mai < perl > jaisa kya hai ?	the homonym for perl is pearl	hom
< council > jaisa kya sunai deta hai ?	the homonym for council is conceal	hom
< city > ko ek vakya mai kaise likhen ?	a sentence with city is: the city takes on an even more interesting hue during event	sent
< fly > ko ek vakya mai kaise likhen ?	a sentence with fly is: airplane fly into a storm cloud	sent
< critique > kai samaan kya hota hai ?	the synonym for critique is evaluate	syn
< psychiatric > kai samaan kya hota hai ?	the synonym for psychiatric is mental	syn

Table 9: Sample x-y pairs in Hindi. The same type of question can be asked in multiple ways.

Question (x)	Answer (y)	type
< edit > de ult ki hunda ae ?	the antonym for edit is expand	ant
< spring > ton bhin ki ae ?	the antonym for spring is compress	ant
< patricide > di paribhasha dasso ?	the definition of patricide is The crime of one who murders his father. Same as Parricide.	defn
< sankhya > da matlab ki hunda ae ?	the definition of sankhya is A Hindoo system of philosophy which refers all things to soul and a rootless germ called prakriti	defn
sunnan vich < site > varga ki ae ?	the homonym for site is sites	hom
< chance > da samnam ki ae ?	the homonym for chance is chants	hom
< enter > nu ek vak vich kidan vartiye ?	a sentence with enter is: as you enter you will be greeted by these rows of trees .	sent
< tree > da prayog ki ae ?	a sentence with tree is: wide shot of an eagle and its nest on top of a tree	sent
< unlawful > jidan ki hunda ae ?	the synonym for unlawful is illegal	syn
< constrict > jidan ki hunda ae ?	the synonym for constrict is close	syn

Table 10: Sample x-y pairs in Punjabi. The same type of question can be asked in multiple ways.

Question (x)	Expected answer (y *)	GPT-3 generated answer	MEM-PROMPT generated answer
give me something i would hear as < whether > ?	the homonym for whether is whither	the example of something you would hear as whether is a question word used to ask about possibilities.	the homonym for whether is weather
what has a < mane > like ring to it ?	the homonym for mane is made	the word that has a mane like ring to it is mane	the homonym for mane is main
What is like < confine > ?	the synonym for confine is hold	the homonym for confine is confined	the synonym for confine is restrain
< leper > means what ?	the definition of leper is a person affected with leprosy.	a sentence with leper is: a leper is a person with a disease that makes his skin scaly and white.	the definition of leper is a person who has leprosy.
expand on < chaperon > ?	the definition of chaperon is a matron who accompanies a young lady in public	the expansion of chaperon is one who chaperons.	the definition of chaperon is a woman who goes with a young girl to social functions.
what is a sentence that can be used to define < wassail > ?	the definition of wassail is an ancient expression of good wishes on a festive occasion	a sentence with wassail is: < wassail > wassails the christmas season	the definition of wassail is to drink to the health of.
expand on < paralbumin > ?	the definition of paralbumin is a proteidlike body found in the fluid from ovarian cysts and elsewhere. it is generally associated with a substance related to	the expansion of paralbumin is an albuminoid globulin occurring in the blood of the horse and other animals, and in the urine of humans.	the definition of paralbumin is a particular kind of albumin.
what can one confuse with < kew > ?	the homonym for kew is key	a word that can be confused with kew is due	the homonym for kew is cue
What is like < pneumatic > ?	the synonym for pneumatic is windy	the homonym for pneumatic is hydraulic	the synonym for pneumatic is airy
give me something i would hear as < seeder > ?	the homonym for seeder is seer	something you would hear as seeder is: someone who plants things in the ground or causes seeds to be planted.	the homonym for seeder is said

Table 11: A subset of random samples where GPT-3-175B without memory was incorrect. The complete set of outputs is located in the anonymized repository <https://github.com/madaan/memprompt/results/results.csv>

Domain	% Finished	Question	Neighbor 1	Neighbor 2	Neighbor 3
Alma mater	1	what high-school did harper lee go to?	what did st augustine do?	who is keyshia cole dad?	when did charles goodyear invent rubber?
Alma mater	5	what college did albert einstein go to?	what high-school did harper lee go to?	who did tim tebow play college football for?	what timezone is utah in?
Alma mater	10	what university did gordon brown attend?	what all does google now do?	what team did david beckham play for in 2011?	who did tim tebow play college football for?
Alma mater	40	where did john mayer go to college?	what school did michael jackson go to high school?	where did derek fisher go to college?	what style of music does john mayer play?
Alma mater	75	where did john steinbeck go to college?	where did john mayer go to college?	what college did john stockton go to?	where did otto frank go to college?
Alma mater	95	where did scott fitzgerald go to college?	who was f. scott fitzgerald?	where did otto frank go to college?	where did derek fisher go to college?
Soccer	1	what team did david beckham play for in 2011?	who did tim tebow play college football for?	what super bowl did peyton manning win?	what type of music did john lennon sing?
Soccer	25	what team did ronaldo play for in 2003?	what part did winona ryder play in star trek?	what to do in richardson dallas?	who did the voice of darth vader in episode 3?
Soccer	33	who did nasri play for before arsenal?	what year did ray allen join the nba?	who does donnie wahlberg play in the sixth sense?	what does david beckham play?
Soccer	65	who has pudge rodriguez played for?	who does nolan ryan play for?	who did carlos boozer play for?	who does ronaldinho play for now 2011?
Soccer	99	what team did david beckham play for before la galaxy?	who does david beckham play for?	what does david beckham play?	what team does david beckham play for in 2012?

Table 12: Relevant examples

Domain	% Finished	Question	Neighbor 1	Neighbor 2	Neighbor 3
Language	1	what does ja- maican people speak?	when was an- cient egypt cre- ated?	where is the denver bron- cos stadium located?	what is the name of the capital of spain?
Language	20	what are the two official languages of paraguay?	what do por- tuguese people speak?	what language does cuba speak?	where is mission san buenaventura located?
Language	37	what language does colombia?	what language does cuba speak?	what was the first language spoken in spain?	what is ser- bian language called?
Language	85	what language does peru speak?	what are the of- ficial languages of the eu?	where is the latin language from?	what do por- tuguese people speak?
Language	90	what language do they speak in colombia south america?	how many lan- guages do they speak in spain?	where is the latin language from?	what language does cuba speak?

Table 13: Relevant examples